

```
In [1]: import numpy as np
import pandas as pd

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

from sklearn.linear_model import LinearRegression, LogisticRegression
from sklearn.neural_network import MLPClassifier, MLPRegressor

from sklearn.tree import DecisionTreeRegressor as dtr
from sklearn.ensemble import RandomForestRegressor

from sklearn.metrics import mean_squared_error as mse
from sklearn.metrics import mean_absolute_error as mae
from sklearn.metrics import r2_score

import warnings
warnings.filterwarnings(action='ignore')
```

```
In [2]: df=pd.read_csv('C:/Users/haris/Downloads/forestfires.csv')
```

```
In [3]: df
```

```
Out[3]:
```

	X	Y	month	day	FFMC	DMC	DC	ISI	temp	RH	wind	rain	area
0	7	5	mar	fri	86.2	26.2	94.3	5.1	8.2	51	6.7	0.0	0.0
1	7	4	oct	tue	90.6	35.4	669.1	6.7	18.0	33	0.9	0.0	0.0
2	7	4	oct	sat	90.6	43.7	686.9	6.7	14.6	33	1.3	0.0	0.0
3	8	6	mar	fri	91.7	33.3	77.5	9.0	8.3	97	4.0	0.2	0.0
4	8	6	mar	sun	89.3	51.3	102.2	9.6	11.4	99	1.8	0.0	0.0
...	...	...	...	...	...	...	...	...	...	...	...	...	...
512	4	3	aug	sun	81.6	56.7	665.6	1.9	27.8	32	2.7	0.0	6.4
513	2	4	aug	sun	81.6	56.7	665.6	1.9	21.9	71	5.8	0.0	54.2
514	7	4	aug	sun	81.6	56.7	665.6	1.9	21.2	70	6.7	0.0	11.1
515	1	4	aug	sat	94.4	146.0	614.7	11.3	25.6	42	4.0	0.0	0.0
516	6	3	nov	tue	79.5	3.0	106.7	1.1	11.8	31	4.5	0.0	0.0

517 rows × 13 columns

```
In [4]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 517 entries, 0 to 516
Data columns (total 13 columns):
#   Column      Non-Null Count  Dtype
---  -
0    X           517 non-null    int64
1    Y           517 non-null    int64
2    month       517 non-null    object
3    day         517 non-null    object
4    FFMC        517 non-null    float64
5    DMC         517 non-null    float64
6    DC          517 non-null    float64
7    ISI         517 non-null    float64
8    temp        517 non-null    float64
9    RH          517 non-null    int64
10   wind        517 non-null    float64
11   rain        517 non-null    float64
12   area        517 non-null    float64
dtypes: float64(8), int64(3), object(2)
memory usage: 52.6+ KB

```

```
In [5]: df['month'].unique()
```

```
Out[5]: array(['mar', 'oct', 'aug', 'sep', 'apr', 'jun', 'jul', 'feb', 'jan',
              'dec', 'may', 'nov'], dtype=object)
```

```
In [6]: df['day'].unique()
```

```
Out[6]: array(['fri', 'tue', 'sat', 'sun', 'mon', 'wed', 'thu'], dtype=object)
```

```
In [7]: df.head(10)
```

```
Out[7]:
```

	X	Y	month	day	FFMC	DMC	DC	ISI	temp	RH	wind	rain	area
0	7	5	mar	fri	86.2	26.2	94.3	5.1	8.2	51	6.7	0.0	0.0
1	7	4	oct	tue	90.6	35.4	669.1	6.7	18.0	33	0.9	0.0	0.0
2	7	4	oct	sat	90.6	43.7	686.9	6.7	14.6	33	1.3	0.0	0.0
3	8	6	mar	fri	91.7	33.3	77.5	9.0	8.3	97	4.0	0.2	0.0
4	8	6	mar	sun	89.3	51.3	102.2	9.6	11.4	99	1.8	0.0	0.0
5	8	6	aug	sun	92.3	85.3	488.0	14.7	22.2	29	5.4	0.0	0.0
6	8	6	aug	mon	92.3	88.9	495.6	8.5	24.1	27	3.1	0.0	0.0
7	8	6	aug	mon	91.5	145.4	608.2	10.7	8.0	86	2.2	0.0	0.0
8	8	6	sep	tue	91.0	129.5	692.6	7.0	13.1	63	5.4	0.0	0.0
9	7	5	sep	sat	92.5	88.0	698.6	7.1	22.8	40	4.0	0.0	0.0

## Preprocessing

Text data included with the numeric data(Month & Days). So we need to encode that in some numeric form before splitting the train test data.

```
In [8]: def ordinal_encoding(df,column,ordering):
df=df.copy()
df[column]=df[column].apply(lambda x: ordering.index(x))
return df
```

```
In [9]: def preprocessing(df,task):
df=df.copy()
df=ordinal_encoding(
    df,
    column='month',
    ordering=[
        'jan', 'feb', 'mar', 'apr', 'may', 'jun', 'jul', 'aug', 'se
        'oct', 'nov', 'dec'
    ]
)
df=ordinal_encoding(
    df,
    column='day',
    ordering=['mon', 'tue', 'wed', 'thu', 'fri', 'sat', 'sun']
)
if task=='Regression':
    Y=df['area']
elif task=='Classification':
    Y=df['area'].apply(lambda x: 1 if x>0 else 0)

X=df.drop('area',axis=1)

X_train,X_test,Y_train,Y_test=train_test_split(X,Y,train_size=0.65,shuffle

scaler=StandardScaler()
scaler.fit(X_train)

X_train=pd.DataFrame(scaler.transform(X_train),columns=X.columns)
X_test=pd.DataFrame(scaler.transform(X_test),columns=X.columns)

return X_train,X_test,Y_train,Y_test
```

## Splitting & Testing Models:

```
In [10]: X_train,X_test,Y_train,Y_test=preprocessing(df,task='Regression')
```

```
In [11]: X_train.head()
```

```
Out[11]:
```

	X	Y	month	day	FFMC	DMC	DC	
0	1.460144	1.330887	0.226770	-0.599195	-0.750425	-0.262938	0.268405	-0
1	0.603934	-0.246461	0.226770	-0.109342	0.748288	0.383050	0.145783	0
2	-0.252276	0.542213	-1.913533	0.380511	0.201882	-1.156059	-1.850008	0
3	-1.108486	0.542213	0.226770	-1.089047	0.264329	0.709954	0.461893	1
4	1.460144	1.330887	-1.913533	0.380511	0.201882	-1.116955	-1.836870	-0

```
In [12]: X_test.head()
```

```
Out[12]:
```

	X	Y	month	day	FFMC	DMC	DC
0	-1.108486	-1.823808	0.226770	-1.089047	0.264329	0.709954	0.461893
1	0.603934	0.542213	0.226770	0.870363	-0.032292	-0.161269	0.326531
2	-0.252276	1.330887	0.654831	-0.109342	0.514114	-0.411530	0.569386
3	-0.680381	-0.246461	0.654831	0.380511	0.451668	0.531643	0.683647
4	1.032039	-0.246461	0.654831	-0.599195	-0.047903	-0.380248	0.770438

```
In [13]: Y_train.head()
```

```
Out[13]:
```

171	2.69
161	1.90
69	0.00
272	3.09
91	0.00

Name: area, dtype: float64

```
In [14]: Y_test.head()
```

```
Out[14]:
```

270	0.52
90	0.00
133	0.00
221	35.88
224	37.71

Name: area, dtype: float64

### Linear Regression

```
In [15]: #The maximum val. of R^2 can be 1.0 that signifies that Linear Regression is  
#Here R^2 score is 0.02051 which is really low signifying that LinearRegression  
linear_reg_model=LinearRegression()  
linear_reg_model.fit(X_train,Y_train)
```

```
print("Performance of Linear Regression R^2 metric {:.5f}".format(linear_reg
```

Performance of Linear Regression R^2 metric 0.02051

### MLP Regressor Model

```
In [16]: # We can see the MLP Regressor Model Works better than Linear Regression,how  
mlp_reg_model=MLPRegressor().fit(X_train,Y_train)  
print("Performance of MLP Regressor Model R^2 metric {:.5f}".format(mlp_reg_
```

Performance of MLP Regressor Model R^2 metric 0.06170

### Decision Tree Regressor

```
In [17]: #R2 score negative not a good fit!  
reg = dtr(random_state = 42)  
reg.fit(X_train, Y_train)
```

```
Y_pred = reg.predict(X_test)
print("MSE =", mse(Y_pred, Y_test))
print("MAE =", mae(Y_pred, Y_test))
print("R2 Score =", r2_score(Y_pred, Y_test))
```

```
MSE = 11504.043293370165
MAE = 27.108121546961325
R2 Score = -7.943021164352407
```

### Random Forest Regressor

```
In [18]: #This works even worse than Decision Tree Regressor
regr = RandomForestRegressor(max_depth=2, random_state=0, n_estimators=100)
regr.fit(X_train, Y_train)
Y_pred = regr.predict(X_test)
print("MSE =", mse(Y_pred, Y_test))
print("MAE =", mae(Y_pred, Y_test))
print("R2 Score =", r2_score(Y_pred, Y_test))
```

```
MSE = 9900.080101011921
MAE = 22.488793674759876
R2 Score = -588.768145917561
```

Reaction after seeing the  $R^2$  Score everytime getting less and less!!!!!!

### Logistic Regression

```
In [19]: X_train,X_test,Y_train,Y_test=preprocessing(df,task='Classification')
```

```
In [20]: Y_train
```

```
Out[20]: 171    1
         161    1
         69    0
         272    1
         91    0
         ..
         129    0
         144    1
         72    0
         235    1
         37    0
         Name: area, Length: 336, dtype: int64
```

```
In [21]: log_reg_model=LogisticRegression()
log_reg_model.fit(X_train,Y_train)

print('Logistic Regression Accuracy, {:.5f}%'.format(log_reg_model.score(X_t
```

```
Logistic Regression Accuracy, 53.59116%
```

### Neural Network Classifier

```
In [22]: nn_classifier_model=MLPClassifier(activation='relu',hidden_layer_sizes=(16,1
nn_classifier_model.fit(X_train,Y_train)
```

Out[22]:

MLPClassifier

MLPClassifier(hidden\_layer\_sizes=(16, 16), n\_iter\_no\_change=100)

In [23]:

```
print('MLP Classifier Accuracy, {:.5f}%'.format(nn_classifier_model.score(X_
```

MLP Classifier Accuracy, 55.80110%

This notebook was converted with [convert.ploomber.io](https://convert.ploomber.io)